



Implementing distributed systems with



Matúš Valo

Distributed systems

- Programming distributed systems is **hard**
- Challenges [1]:
 - Service access and configuration
 - Event handling
 - Concurrency
 - Synchronization
- Single point of failure problem
- Split brain problem
- CAP theorem



Failover Celery Beat

- **celery beat** is a scheduler; It kicks off tasks at regular intervals, that are then executed by available worker nodes in the cluster.

The screenshot shows the GitHub interface for the 'celery/celery' repository. At the top, the repository name 'celery / celery' is displayed. To the right, there are buttons for 'Watch' (398), 'Star' (8,824), and 'Fork' (2,521). Below these, a navigation bar includes links for 'Code', 'Issues' (273), 'Pull requests' (38), 'Projects' (0), 'Wiki', and 'Insights'. The 'Issues' section is active, showing a specific issue titled 'Failover for celerybeat #1495'. The issue is marked as 'Open' and was opened by 'kerwin' on 'Aug 1, 2013' with '23 comments'. A green 'New issue' button is visible in the top right corner of the issue view.

celery / celery

Watch 398 Star 8,824 Fork 2,521

<> Code Issues 273 Pull requests 38 Projects 0 Wiki Insights

Failover for celerybeat #1495

Open kerwin opened this issue on Aug 1, 2013 · 23 comments

New issue



Highly Available Airflow

- Airflow is a platform to programmatically author, schedule and monitor workflows.
- Workflows are Directed acyclic graphs (DAGs) of tasks
- Scheduler executes tasks on an array of workers while following the specified dependencies.
- Currently, Airflow has single point of failure – Scheduler [1]
- Multiple attempts exists for solving this issue including home-brew solutions [2]

[1] <http://site.clairvoyantsoft.com/making-apache-airflow-highly-available/>

[2] <https://github.com/teamclairvoyant/airflow-scheduler-failover-controller>





Why Consul?

- **Consul is clustered and highly available.**
- **Consul enables easy development of distributed and highly available systems.**
- Consul makes it simple for services to register themselves and to discover this services via a DNS or HTTP interface.
- Pairing service discovery with health checking prevents routing requests to unhealthy hosts and enables services to easily provide circuit breakers.
- Consul scales to multiple datacenters out of the box with no complicated configuration. Look up services in other datacenters, or keep the request local.
- Flexible key/value store for dynamic configuration, feature flagging, coordination, leader election and more.

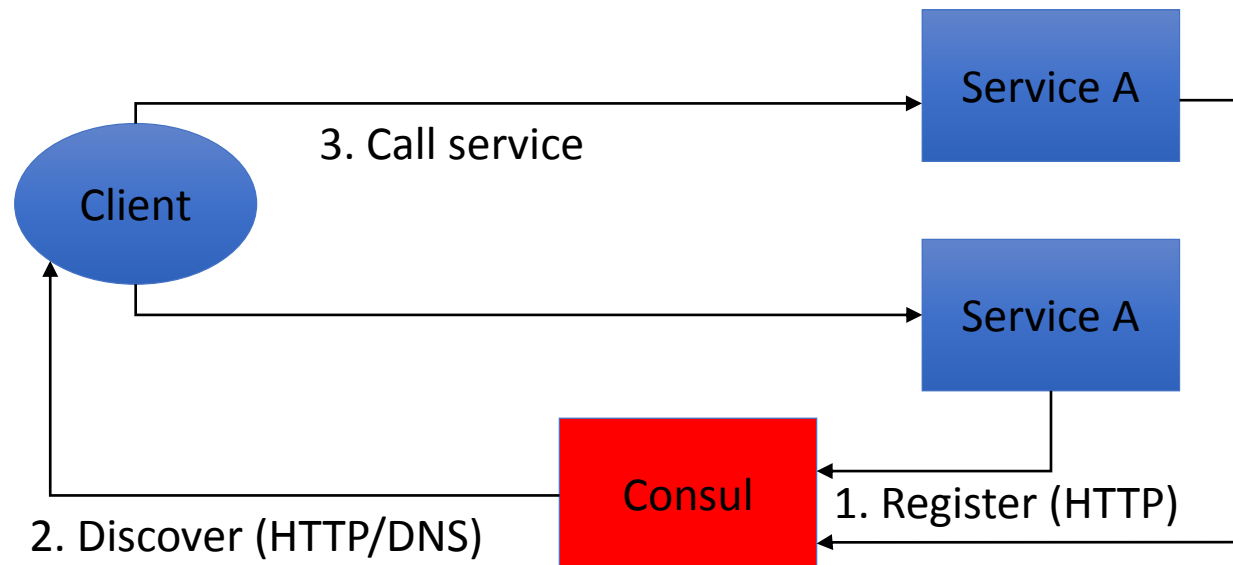


Consul Cluster

- Consul cluster consists from at least 3 nodes (supports 1 node failure)
- When higher failure tolerance is needed:
 - 5 nodes (supports 2 node failure)
 - 7 nodes (supports 3 node failure)
- When new cluster is created it needs to be bootstrapped:
 - each node is started in bootstrap mode using `-bootstrap-expect BNODES` parameter. BNODES is the initial number of nodes in cluster.
 - nodes are joined together using `consul join` command
- consul supports special single node DEV mode for developing purposes.

Service Discovery

- Useful for scalable distributed system, where new nodes needs to be discovered automatically.
- Distributed component can register at startup and from that point is visible



- Service Registration

```
$ cat servicea.json
```

```
{  
  "ID": "servicea1",  
  "Name": "serviceA",  
  "Tags": ["primary", "v1"],  
  "Address": "servicea.example.com",  
  "Port": 8000  
}
```

```
$ curl -X PUT -d @servicea.json localhost:8500/v1/agent/service/register
```

- Service discovery (DNS Interface)

```
$ dig -p 8600 servicea.service.consul @t3dredis01
```

```
/// OUTPUT OMMITED
```

```
;; ANSWER SECTION:
```

```
servicea.service.consul. 0      IN      CNAME   servicea.example.com.
```

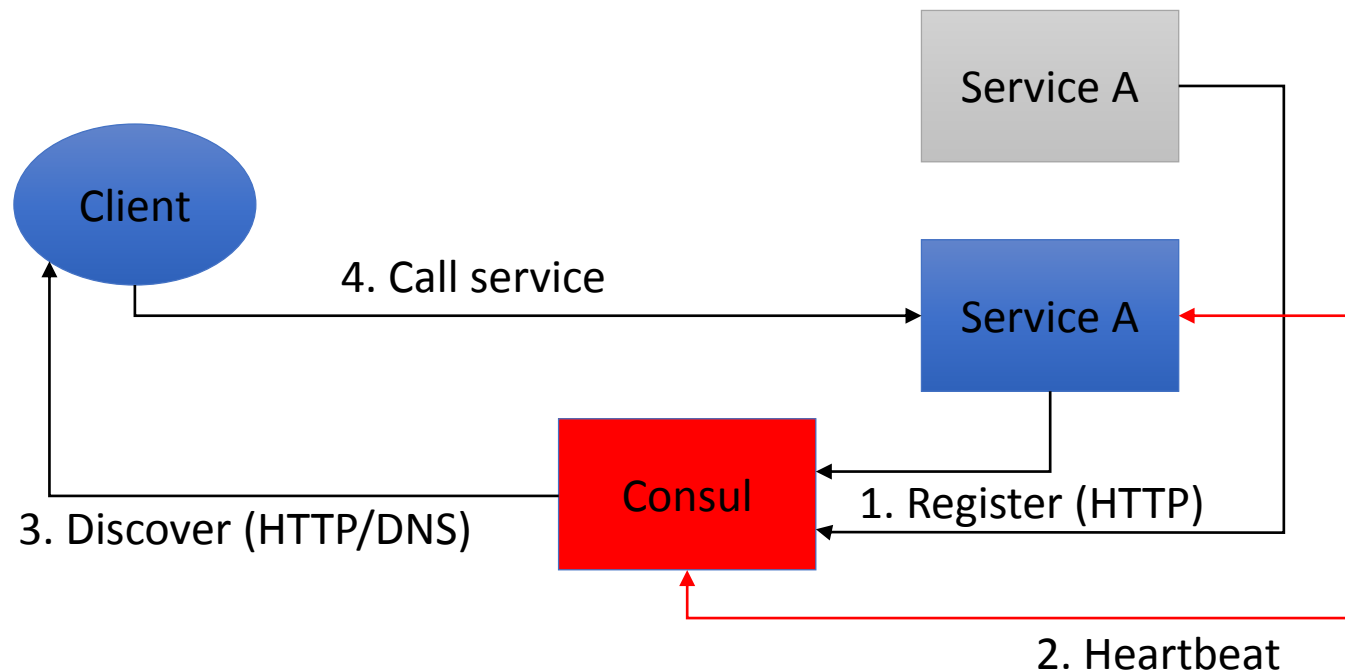
```
/// OUTPUT OMMITED
```

```
$ dig -p 8600 servicea.service.consul @t3dredis01 SRV
```



Health check

- Service can ask consul to track availability of service using heartbeat (HB)
- Multiple ways of HB are supported: HTTP, raw TCP/IP, Custom Script, Docker container check



- Service registration

```
$ cat servicea.json
```

```
{  
  "ID": "servicea1",  
  "Name": "serviceA",  
  "Tags": ["primary", "v1"],  
  "Address": "servicea.example.com",  
  "Port": 8000,  
  "Check": {  
    "DeregisterCriticalServiceAfter": "90m",  
    "HTTP": "http://servicea.example.com:5000/health",  
    "Interval": "10s"  
  }  
}
```

```
$ curl -X PUT -d @servicea.json localhost:8500/v1/agent/service/register
```

- Service discovery (HTTP interface)

```
$ curl localhost:8500/v1/health/service/servicea?passing=true
```

```
[
  {
    "Node": {
      "ID": "3114d0f7-ff8b-fcc9-a8b8-82fde8cc6a4f",
      "Node": "mynode.example.com",
      "Address": "192.168.10.10",
      "Datacenter": "dc1",
      /// OUTPUT OMITTED
    },
    "Service": {
      "ID": "servicea1",
      "Service": "serviceA",
      "Tags": ["primary", "v1"],
      "Address": "servicea.example.com",
      "Port": 8000,
      /// OUTPUT OMITTED
    },
    "Checks": [ /// OUTPUT OMITTED ]
  }
]
```



K/V Storage

- Consul provides Key/Value Storage which is distributed and highly available.
- Any data can be stored in form “key”: “value”
 - “value” can store any JSON value encoded as base64
- K/V Storage has CLI or HTTP API (easy to use)
- Consul also provides additional functionality:
 - transactions – multiple operations on K/V storage with atomic execution.
 - atomic key updates using a Check-And-Set operation

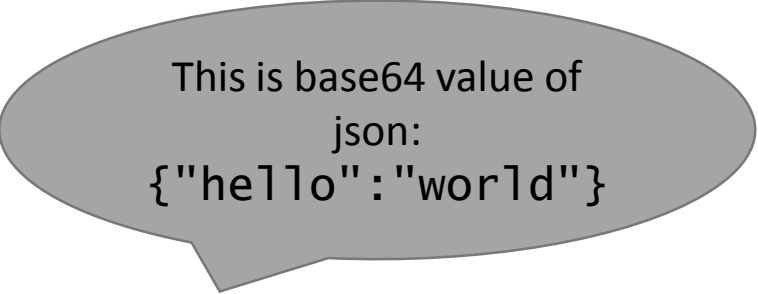
- Setting value to K/V storage

```
$ curl -X PUT -d '{"hello":"world"}' localhost:8500/v1/kv/hello
true
```

- Getting value from K/V storage

```
$ curl localhost:8500/v1/kv/hello
[
```

```
  {
    "LockIndex": 0,
    "Key": "hello",
    "Flags": 0,
    "Value": "eyJ0ZWxsbyI6IndvcmxkIn0=",
    "CreateIndex": 87,
    "ModifyIndex": 87
  }
```



This is base64 value of
json:
{"hello":"world"}


```
]
```

- Deleting value from K/V storage

```
$ curl -X DELETE localhost:8500/v1/kv/hello
true
```

Distributed locks

- Semantic of Consul locks is similar to the mutex (mutual exclusion) but can be used also as a semaphore
- Consul agent must be present on client
- To start process holding lock: `consul lock [options] <LOCK> <CMD>`
- This guarantee that only one single <CMD> is executed in cluster
- When <CMD> exists/crash the lock is automatically freed.



[Unit]

Description=Failover Service

After=consul.service

Requires=consul.service

[Service]

ExecStart=/usr/local/bin/consul lock -verbose <LOCK_NAME> <EXECUTABLE>

Restart=always

KillSignal=SIGQUIT

Type=simple

StandardError=syslog

NotifyAccess=all

[Install]

WantedBy=multi-user.target



Implementing failover Celery Beat

1. Implement custom scheduler with central storage
e.g. DatabaseScheduler from django-celery-beat
2. Start celery beat on at least two nodes using consul distributed lock with custom scheduler

```
$ consul lock celery_beat_lock celery -A proj beat --scheduler  
django_celery_beat.schedulers:DatabaseScheduler
```



Additional Consul features

- Access Control List (ACL) system
- Multi Datacenter
- Commercial support
- Watches and Events
- Web UI
- Libraries for multiple languages (python-consul)



Thank you