

SIMPLIFYING CONTAINERS AT SCALE

---

# THE CONCIERGE PARADIGM

## YOUR CONCIERGE FOR THE EVENING

- ▶ Gareth Brown
- ▶ Director and technologist at Mesoform
- ▶ Specialise in securely simplifying and streamlining
- ▶ DevOps back in early '00s
- ▶ Was running containers in production many years ago
- ▶ Built a self-service VM infrastructure..









## HISTORY OF CONTAINERS

- ▶ 1979: chroot
- ▶ Jails, Zones, LXC (2000, 2004, 2008)
- ▶ Along comes AWS
- ▶ Docked back in



## FLYING FISH

- ▶ Docker Engine in the Cloud
- ▶ Maintaining pets
- ▶ Monitoring
- ▶ Scheduling
- ▶ Auto-scaling
- ▶ Service discovery







## LAYER CAKE

- ▶ New technologies (Kubernetes, Mesos)
- ▶ Complex
- ▶ Integrating different workloads and IaaS
- ▶ Up-skilling and support
- ▶ Tight coupling and dependency
- ▶ Keep It Stupidly Simple



**THEY SHOULD REMAKE "BACK TO THE FUTURE 2"  
WHERE THERE ARE NO FLYING CARS**

**AND PEOPLE JUST STARE AT THEIR PHONES ALL  
THE TIME GETTING OFFENDED AT EVERYTHING**

**SOCIAL DEMENTIA**



## FUTURE OF CONTAINERS

- ▶ Standardisation
- ▶ Portability
- ▶ Performance
- ▶ Simplified management
- ▶ Resource Utilisation
- ▶ Cost!



**PUPPIES MAKE PAIN GO AWAY**



## OPERATING PAINS

- ▶ On-Premise, EC2, ECS, CoreOS, Kubernetes, other AWS services, Java, Python...
- ▶ Papertrail and Elastic Stack
- ▶ Zabbix and Librato
- ▶ Dropwizard with agents pulling from applications
- ▶ bumped all of the common issues
- ▶ Windowing and performance

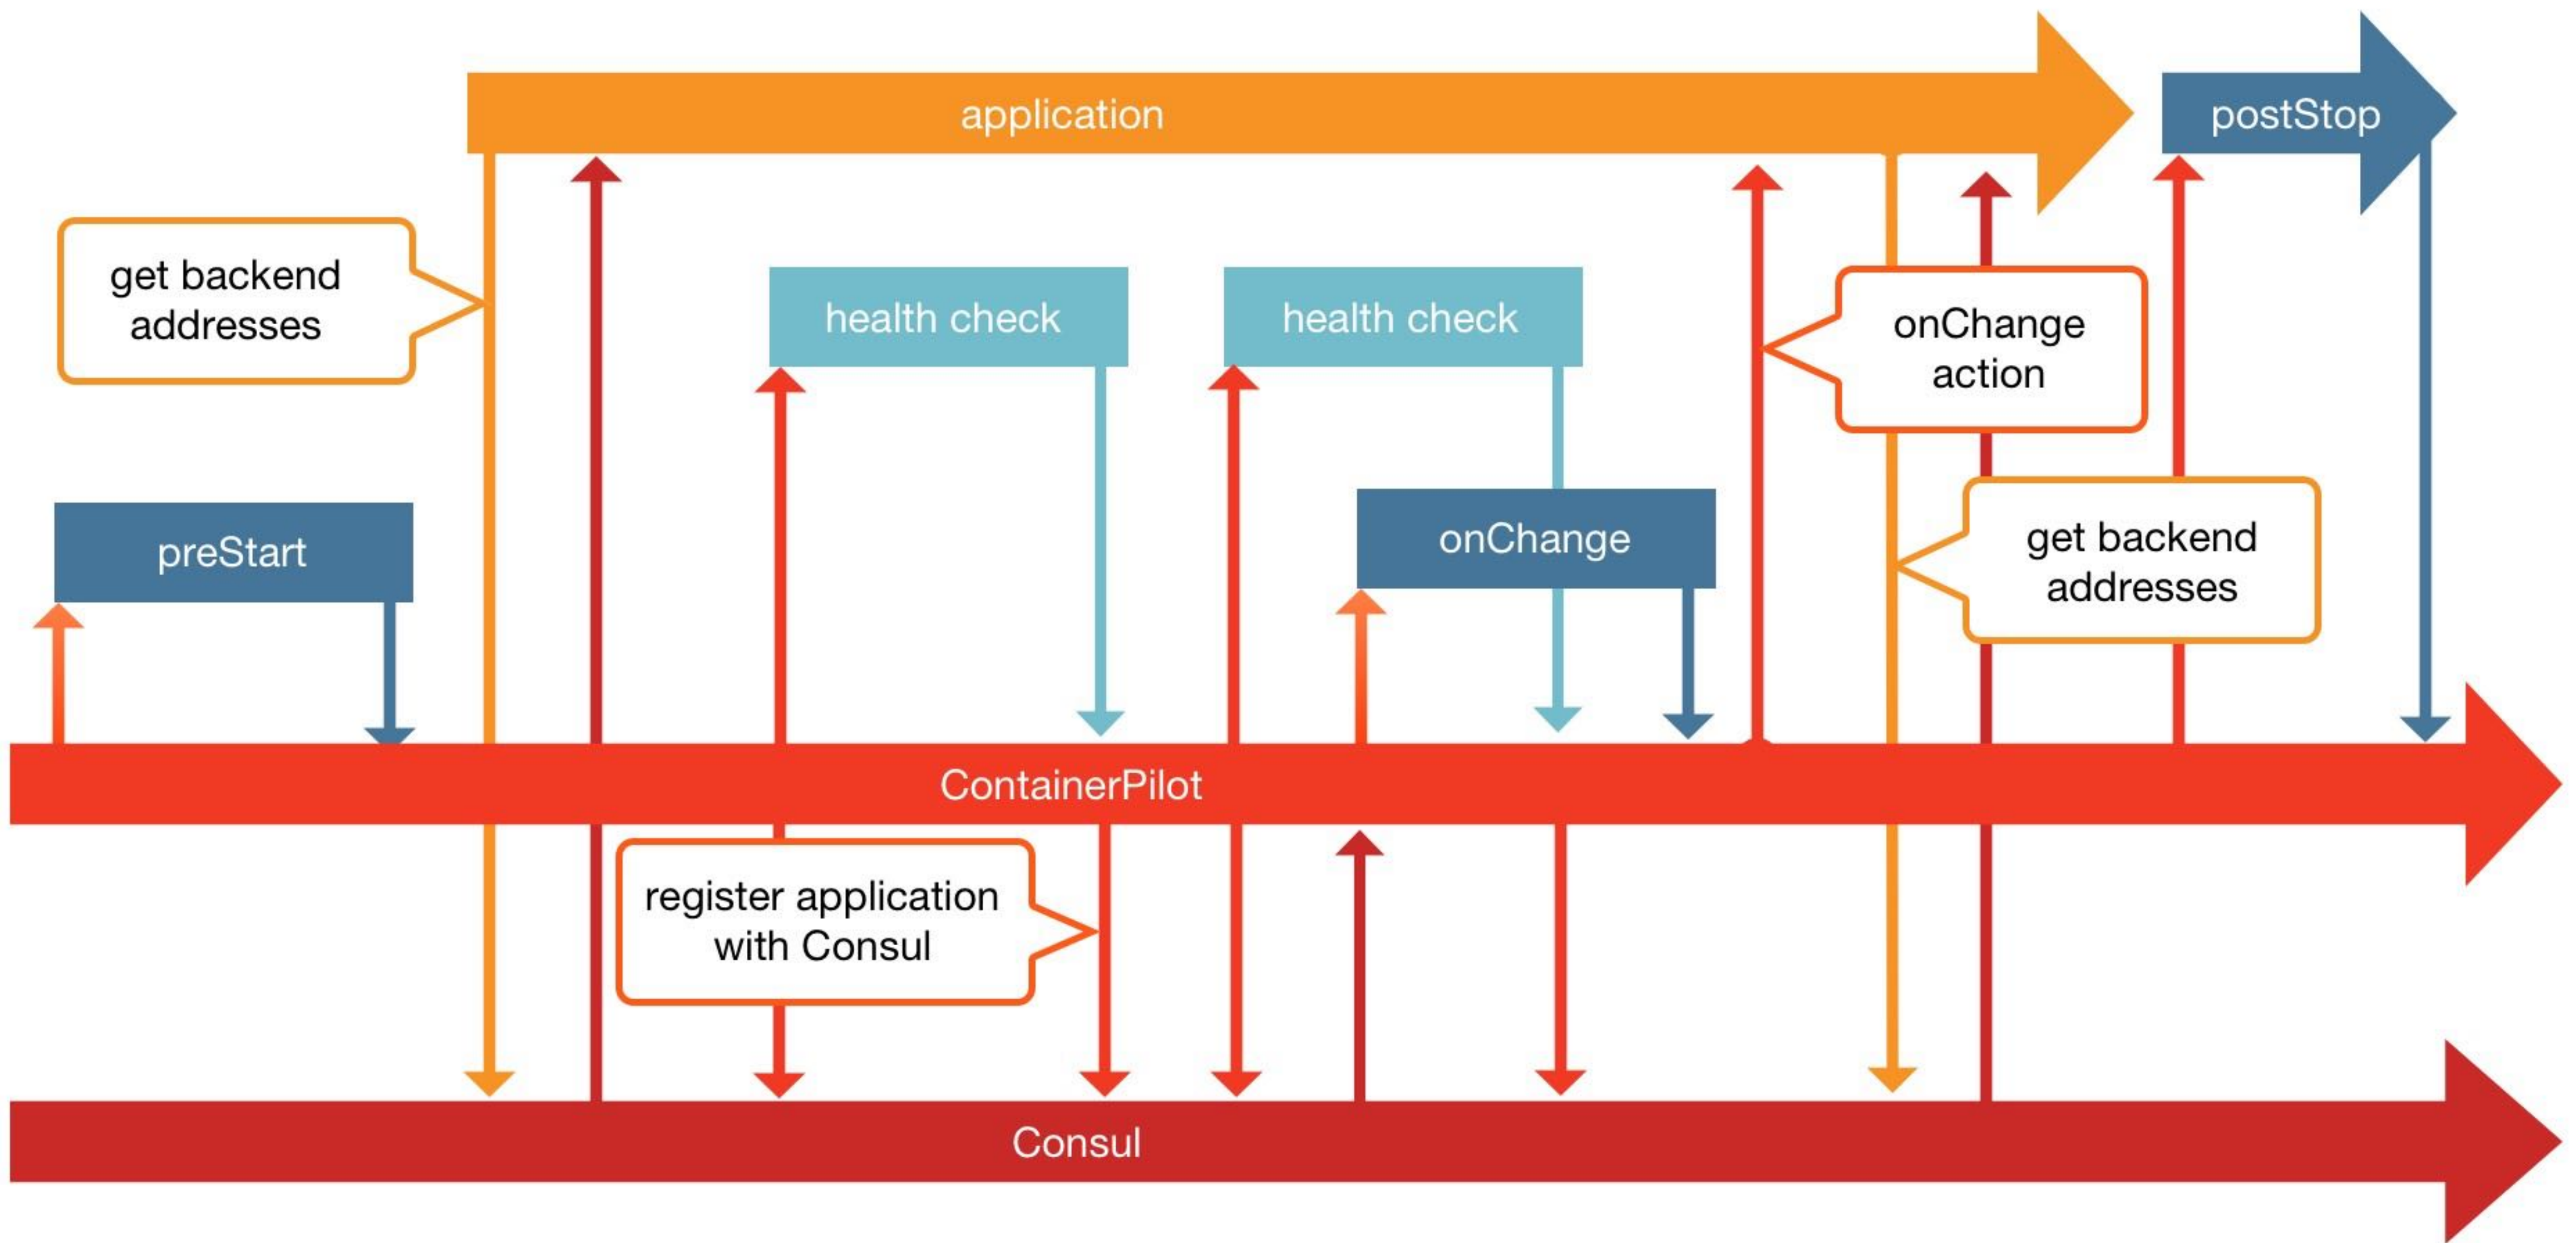






## AUTOPILOT PATTERN

- ▶ No complex framework
- ▶ Service discovery
- ▶ Application orchestration
- ▶ Small piece of code to automate common actions





## FLYING ON AUTOPILOT

- ▶ Scheduler agnostic
- ▶ Most things just work
- ▶ App-centric orchestration
- ▶ Drastically less management
- ▶ Production grade environment, test environment time
- ▶ Co-processes!





CARTOONSTOCK  
.com

Search ID: m5in702

Shiell



## BATTERIES INCLUDED

- ▶ Loose-couple to well defined systems
- ▶ Automatically register our containers
- ▶ Automatically discover resources
- ▶ Self-healing or corrective actions
- ▶ Interact with legacy applications
- ▶ Compliance scanning



## CONTAINERPILOT.JSON

```
jobs: [{
  name: 'scheduling-status-healthy',
  exec: ['zabbix_sender', '-c', '/etc/coproceses/zabbix/zabbix_agentd.conf',
        '--key', 'container.state', '--value', '1'],
  when: { source: 'apache-fwdproxy', each: 'healthy' }
},{
  name: 'zabbix-agent',
  exec: ['/usr/sbin/zabbix_agentd',
        '-fc', '/etc/coproceses/zabbix/zabbix_agentd.conf'],
  restarts: 'unlimited',
  health: { exec: 'zabbix_agentd -t agent.ping', interval: 30, ttl: 60, timeout: 5 },
  when: { source: 'platform-integration-setup', once: 'exitSuccess' }
},{
  name: 'post-stop',
  exec: ['zabbix_sender', '-c', '/etc/coproceses/zabbix/zabbix_agentd.conf',
        '--key', 'container.state', '--value', '0'],
  when: { once: 'shutdown' }
}]
```







## PUSH VS PULL

- ▶ Push method: auto-register but no confidence in instance state
- ▶ Pull method: centralised configuration but extra management
- ▶ Pull understands load and partitioning
- ▶ Processing poor performance
- ▶ Windowing
- ▶ Can we unify push and pull?





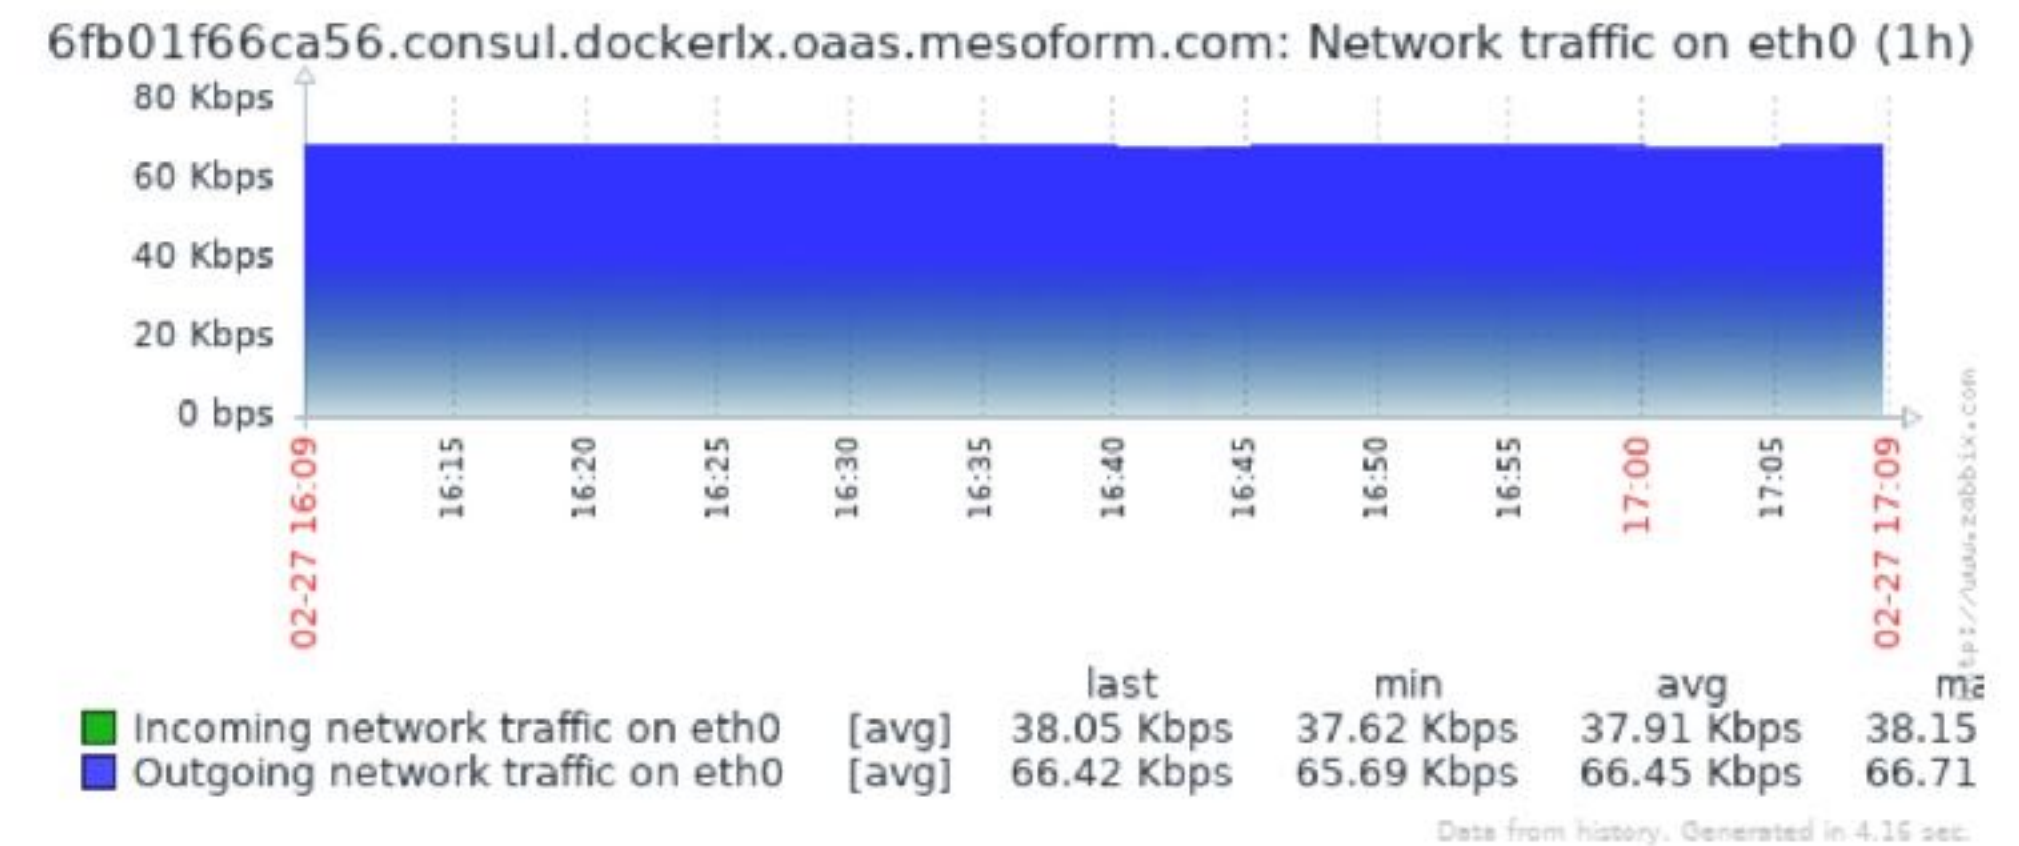







## THE CONCIERGE COURIER

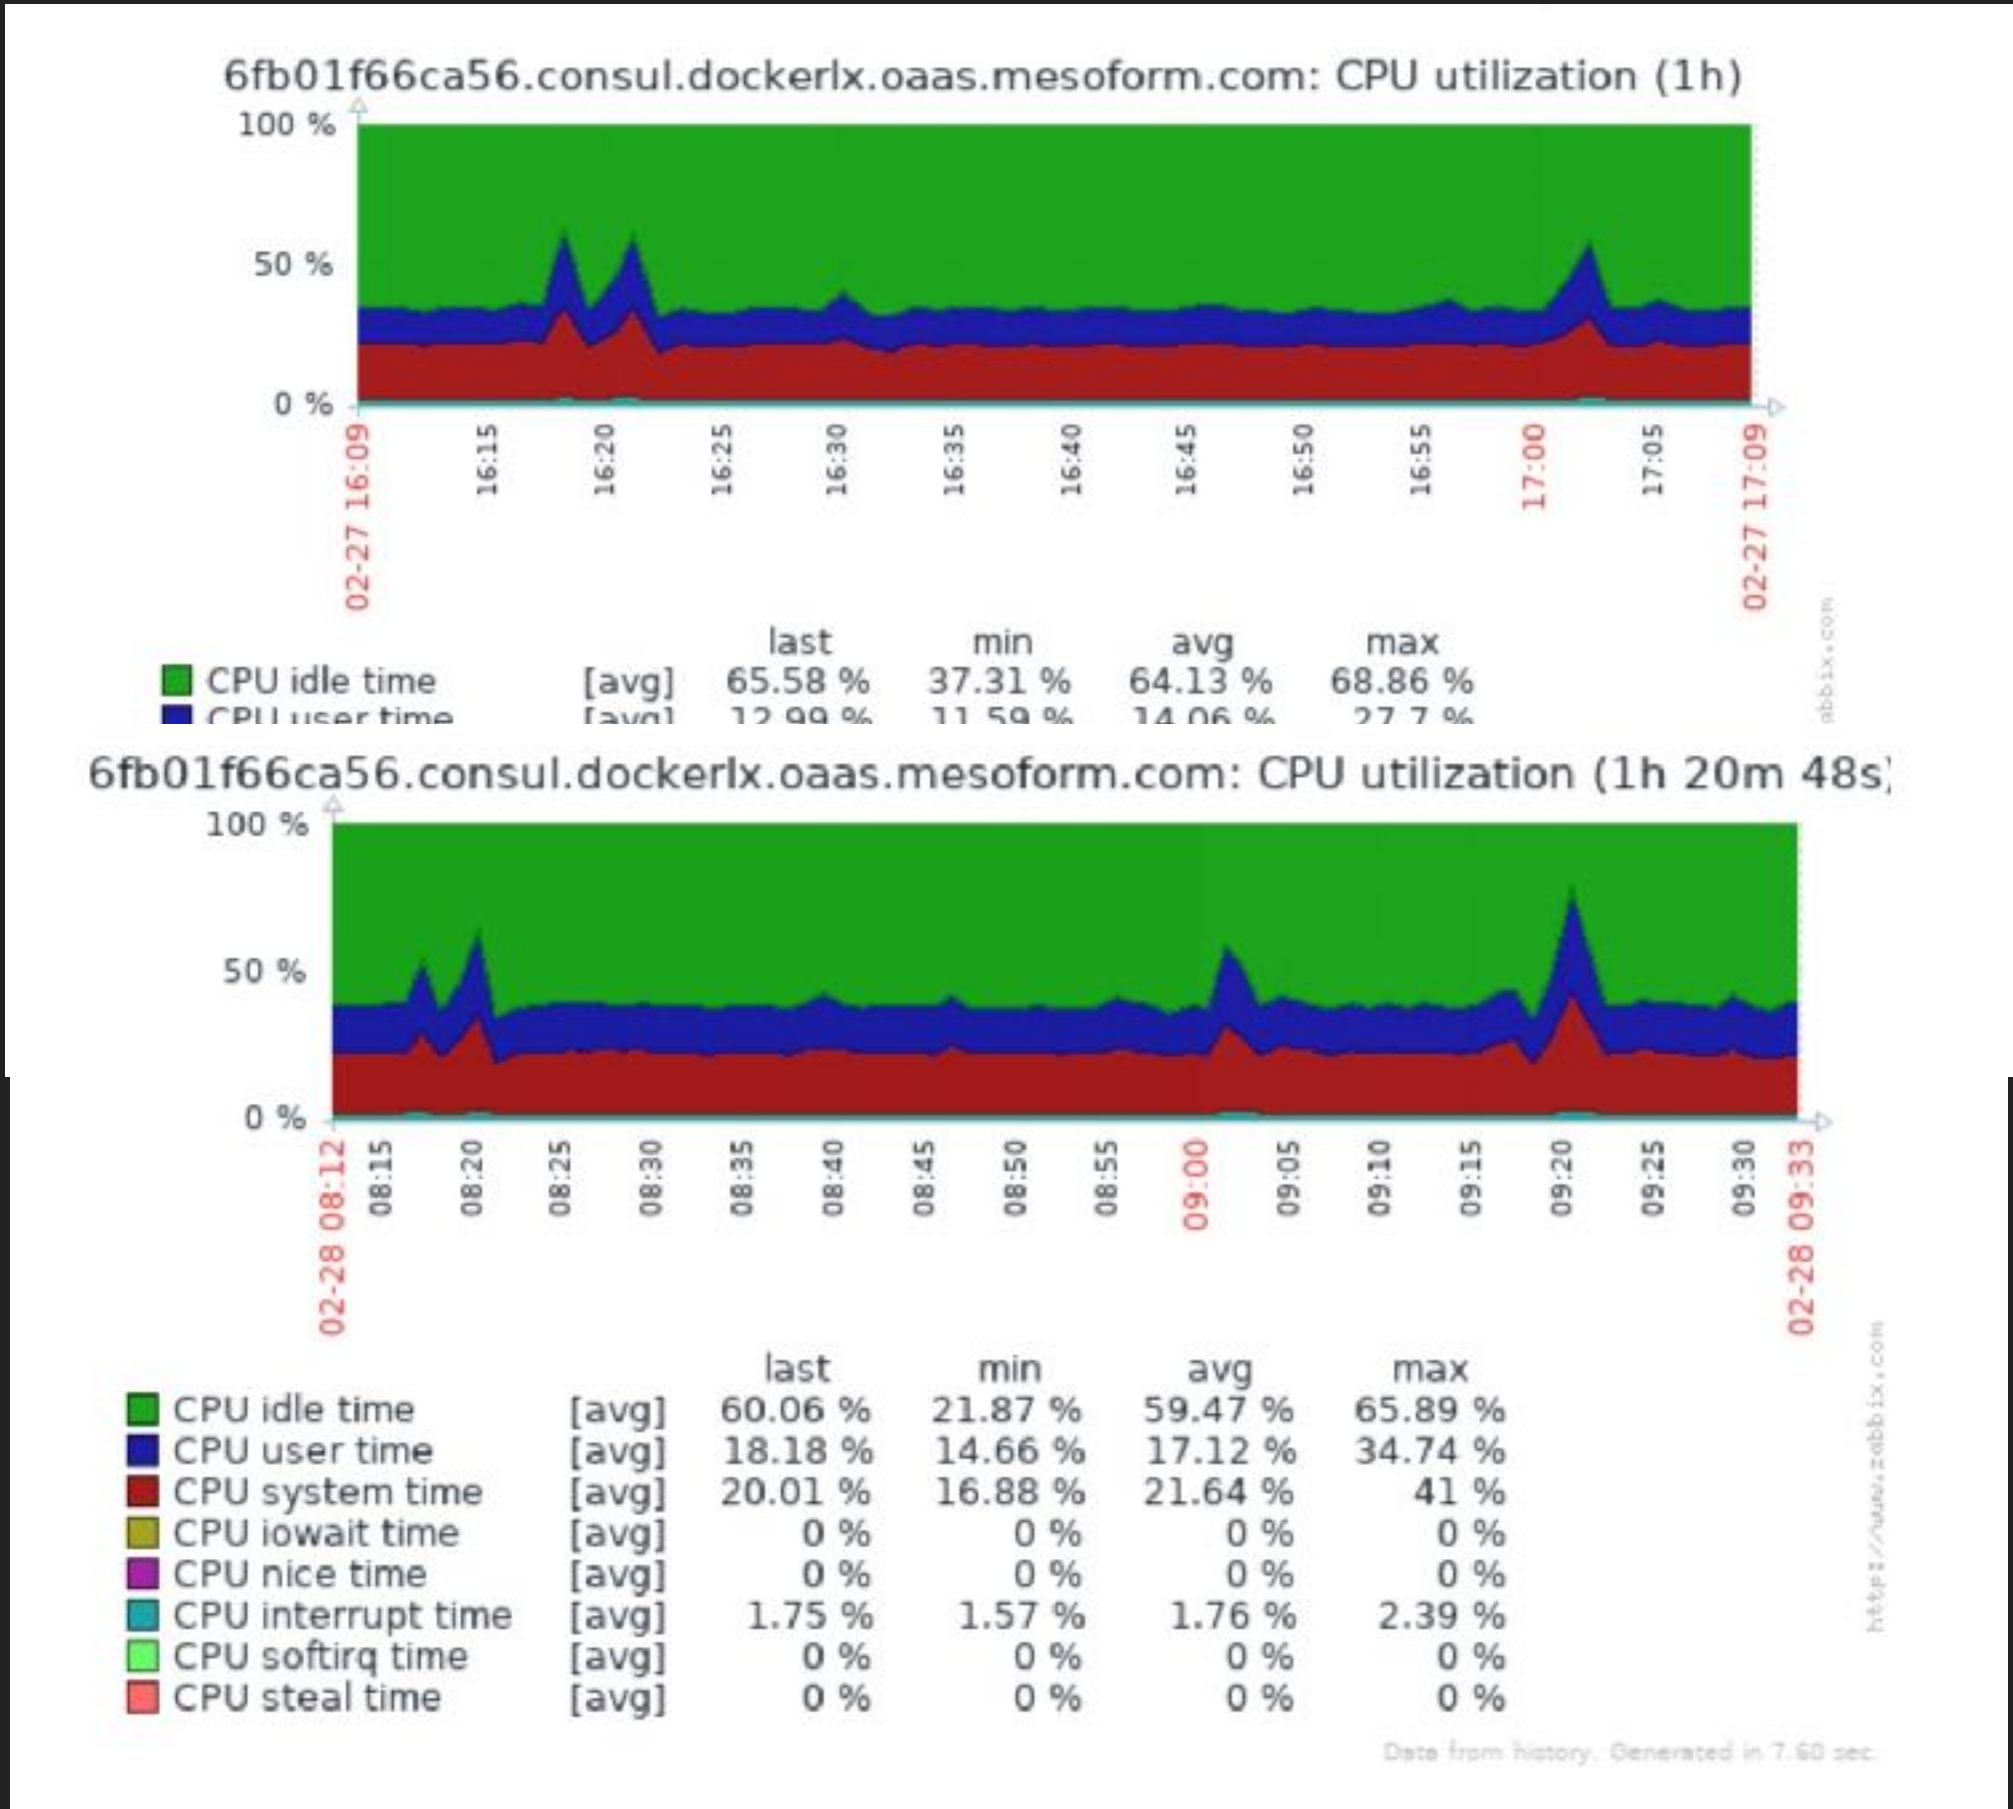
- ▶ Two purposes (discovery, delivery)
- ▶ Learns metrics
- ▶ Picks up metrics
- ▶ Delivers them
- ▶ Records delivery
- ▶ Performance?







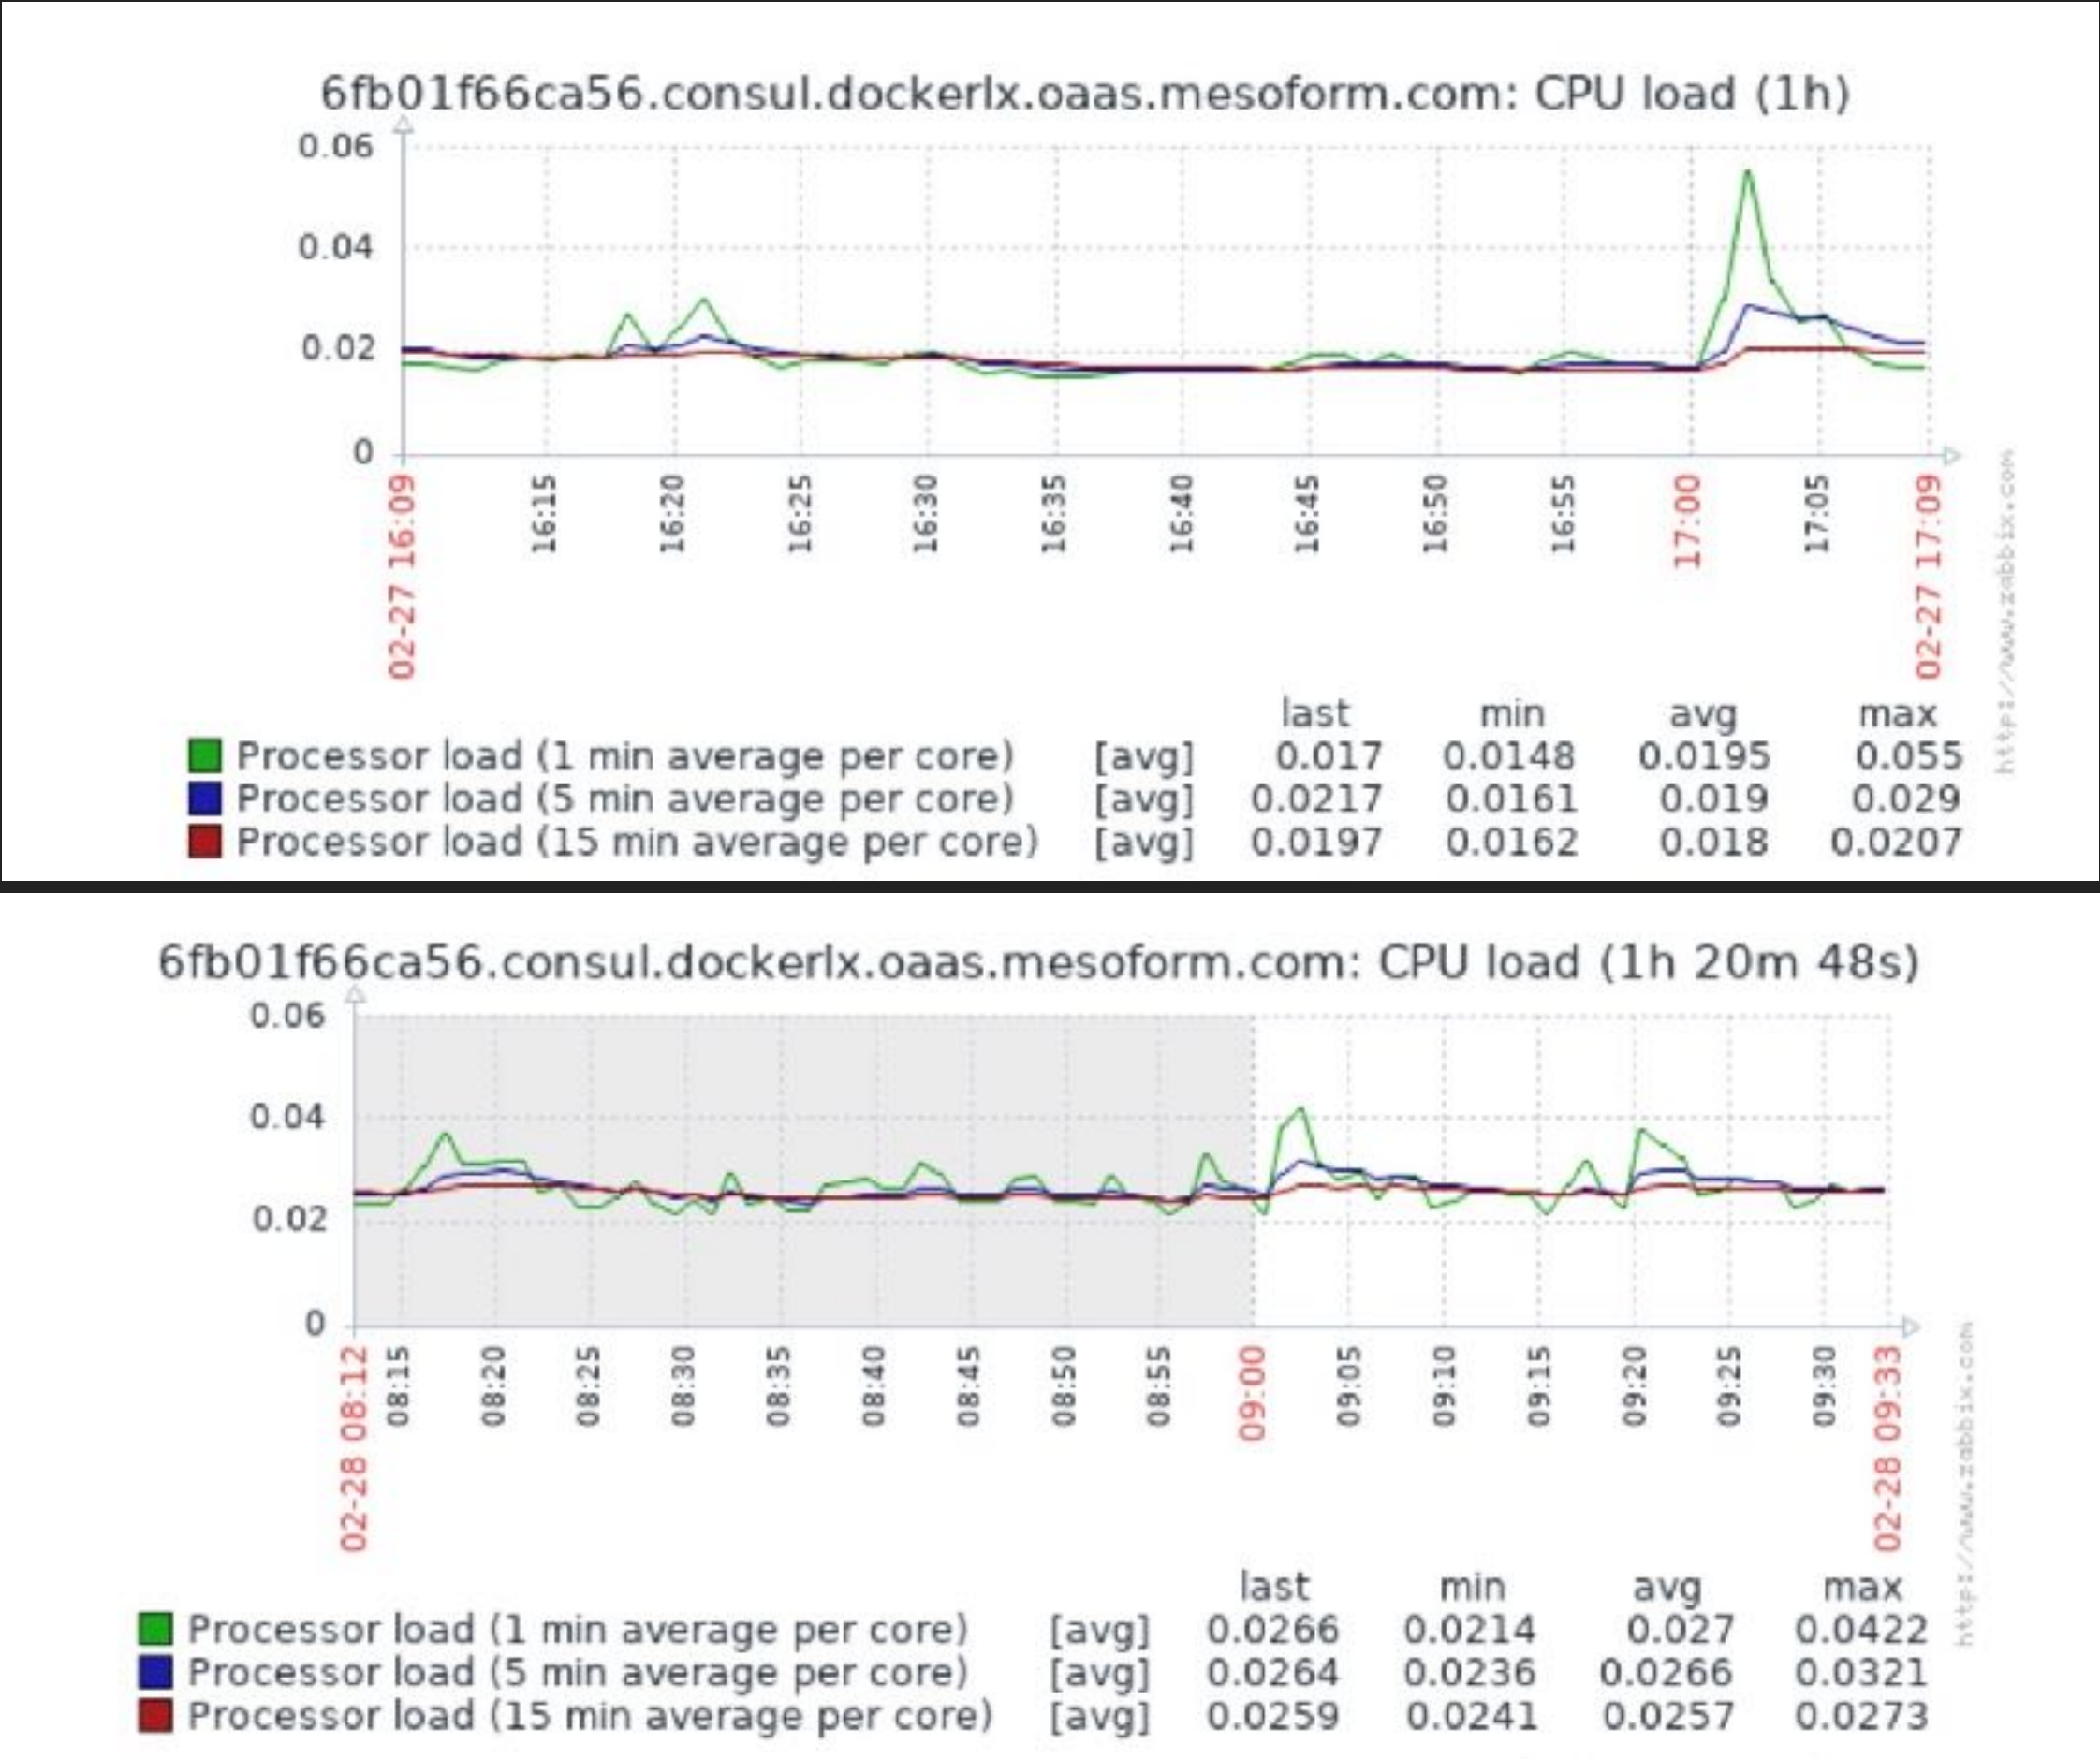
# CPU UTILISATION






# THE CONCIERGE PARADIGM

## LOAD




# THE CONCIERGE PARADIGM

## TIMING







6fb01f66ca56.consul.dockerlx.oaas.mesoform.com: Consul::metrics:collect::timers (1h 20m 48s)




6fb01f66ca56.consul.dockerlx.oaas.mesoform.com: Memory usage (1h 20m 48s)




6fb01f66ca56.consul.dockerlx.oaas.mesoform.com: CPU utilization (1h 20m 48s)




6fb01f66ca56.consul.dockerlx.oaas.mesoform.com: Network traffic on eth0 (1h 20m 48s)



6fb01f66ca56.consul.dockerlx.oaas.mesoform.com: CPU jumps (1h 20m 48s)



6fb01f66ca56.consul.dockerlx.oaas.mesoform.com: CPU load (1h 20m 48s)





## THE CONCIERGE COURIER

- ▶ 3rd party features
- ▶ No windowing
- ▶ High Performance
- ▶ Send to anywhere
- ▶ Pull from anywhere
- ▶ Agnostic



## CONCIERGE\_COURIER.PY

```
def discover_timers():
    """
    Output Zabbix formatted JSON of keys
    """
    # just for testing purposes, simply open a file with metrics
    with open("/tmp/metrics.json", "r") as metrics_file:
        keys = metrics_file.read()
        keys_json = json.loads(keys)

        discovery_data_dict = \
            {'data': [{"{#TIMER}": key} for key in keys_json['timers']]}
    print(json.dumps(discovery_data_dict))
```

## CONCIERGE\_COURIER.PY

```
def get_timers():
    with open("/tmp/metrics.json", "r") as metrics_file:
        keys = metrics_file.read()
        keys = json.loads(keys)
        with open("/tmp/timer_metrics_zabbix.sender", "w") as sender_file:
            for timer_name, metrics in keys['timers'].items():
                for metric_name, metric_value in metrics.items():
                    sender_file.write("- timer[{0}.{1}] {2}\n"
                                      .format(timer_name, metric_name, metric_value))
    send_metrics("timer")

def send_metrics(metric_type):
    filename = "/tmp/" + metric_type + "_metrics_zabbix.sender"
    call("zabbix_sender -c /etc/coproceses/zabbix/zabbix_agentd.conf -i "
        + filename + " >/dev/null", shell=True)
    print time.time() - startTime
```



## THE ENFIELD METHOD

- ▶ Accurate, single-shot, immediate feedback
- ▶ Like the rifle
- ▶ Backoff under network issues
- ▶ Greater confidence in container state
- ▶ Greater confidence in state of whole system
- ▶ More frequent updates

## STATE TO STATE

- ▶ State in service discovery
- ▶ State in event management
- ▶ End-to-end view of whole system
- ▶ State history
- ▶ Dev/Ops on the same page
- ▶ State manipulation!







## STATE CONTROL

- ▶ Consul keeps configuration state
- ▶ Monitoring performance and availability state
- ▶ Dynamic Asset database
- ▶ Automate scheduling, scaling, archiving



**WHAT IF I TOLD YOU**

**SIRI IS A PERSONAL ASSISTANT,  
NOT A TALK-BUDDY**

## THE CONCIERGE SCHEDULER

- ▶ Containers Auto-register
- ▶ Push & pull state
- ▶ Optimised over many years
- ▶ Grouping containers by service
- ▶ Data about whole system
- ▶ Basically just runs *docker-compose scale*



## SCALING

- ▶ Complex trigger profiles
- ▶ Pre-scaling using a predictive trigger
- ▶ Compare upstream service performance as well
- ▶ Vertical scaling
- ▶ Escalation steps
- ▶ Scaling events and problem events in one system

## CONCIERGE\_SCHEDULER.SH

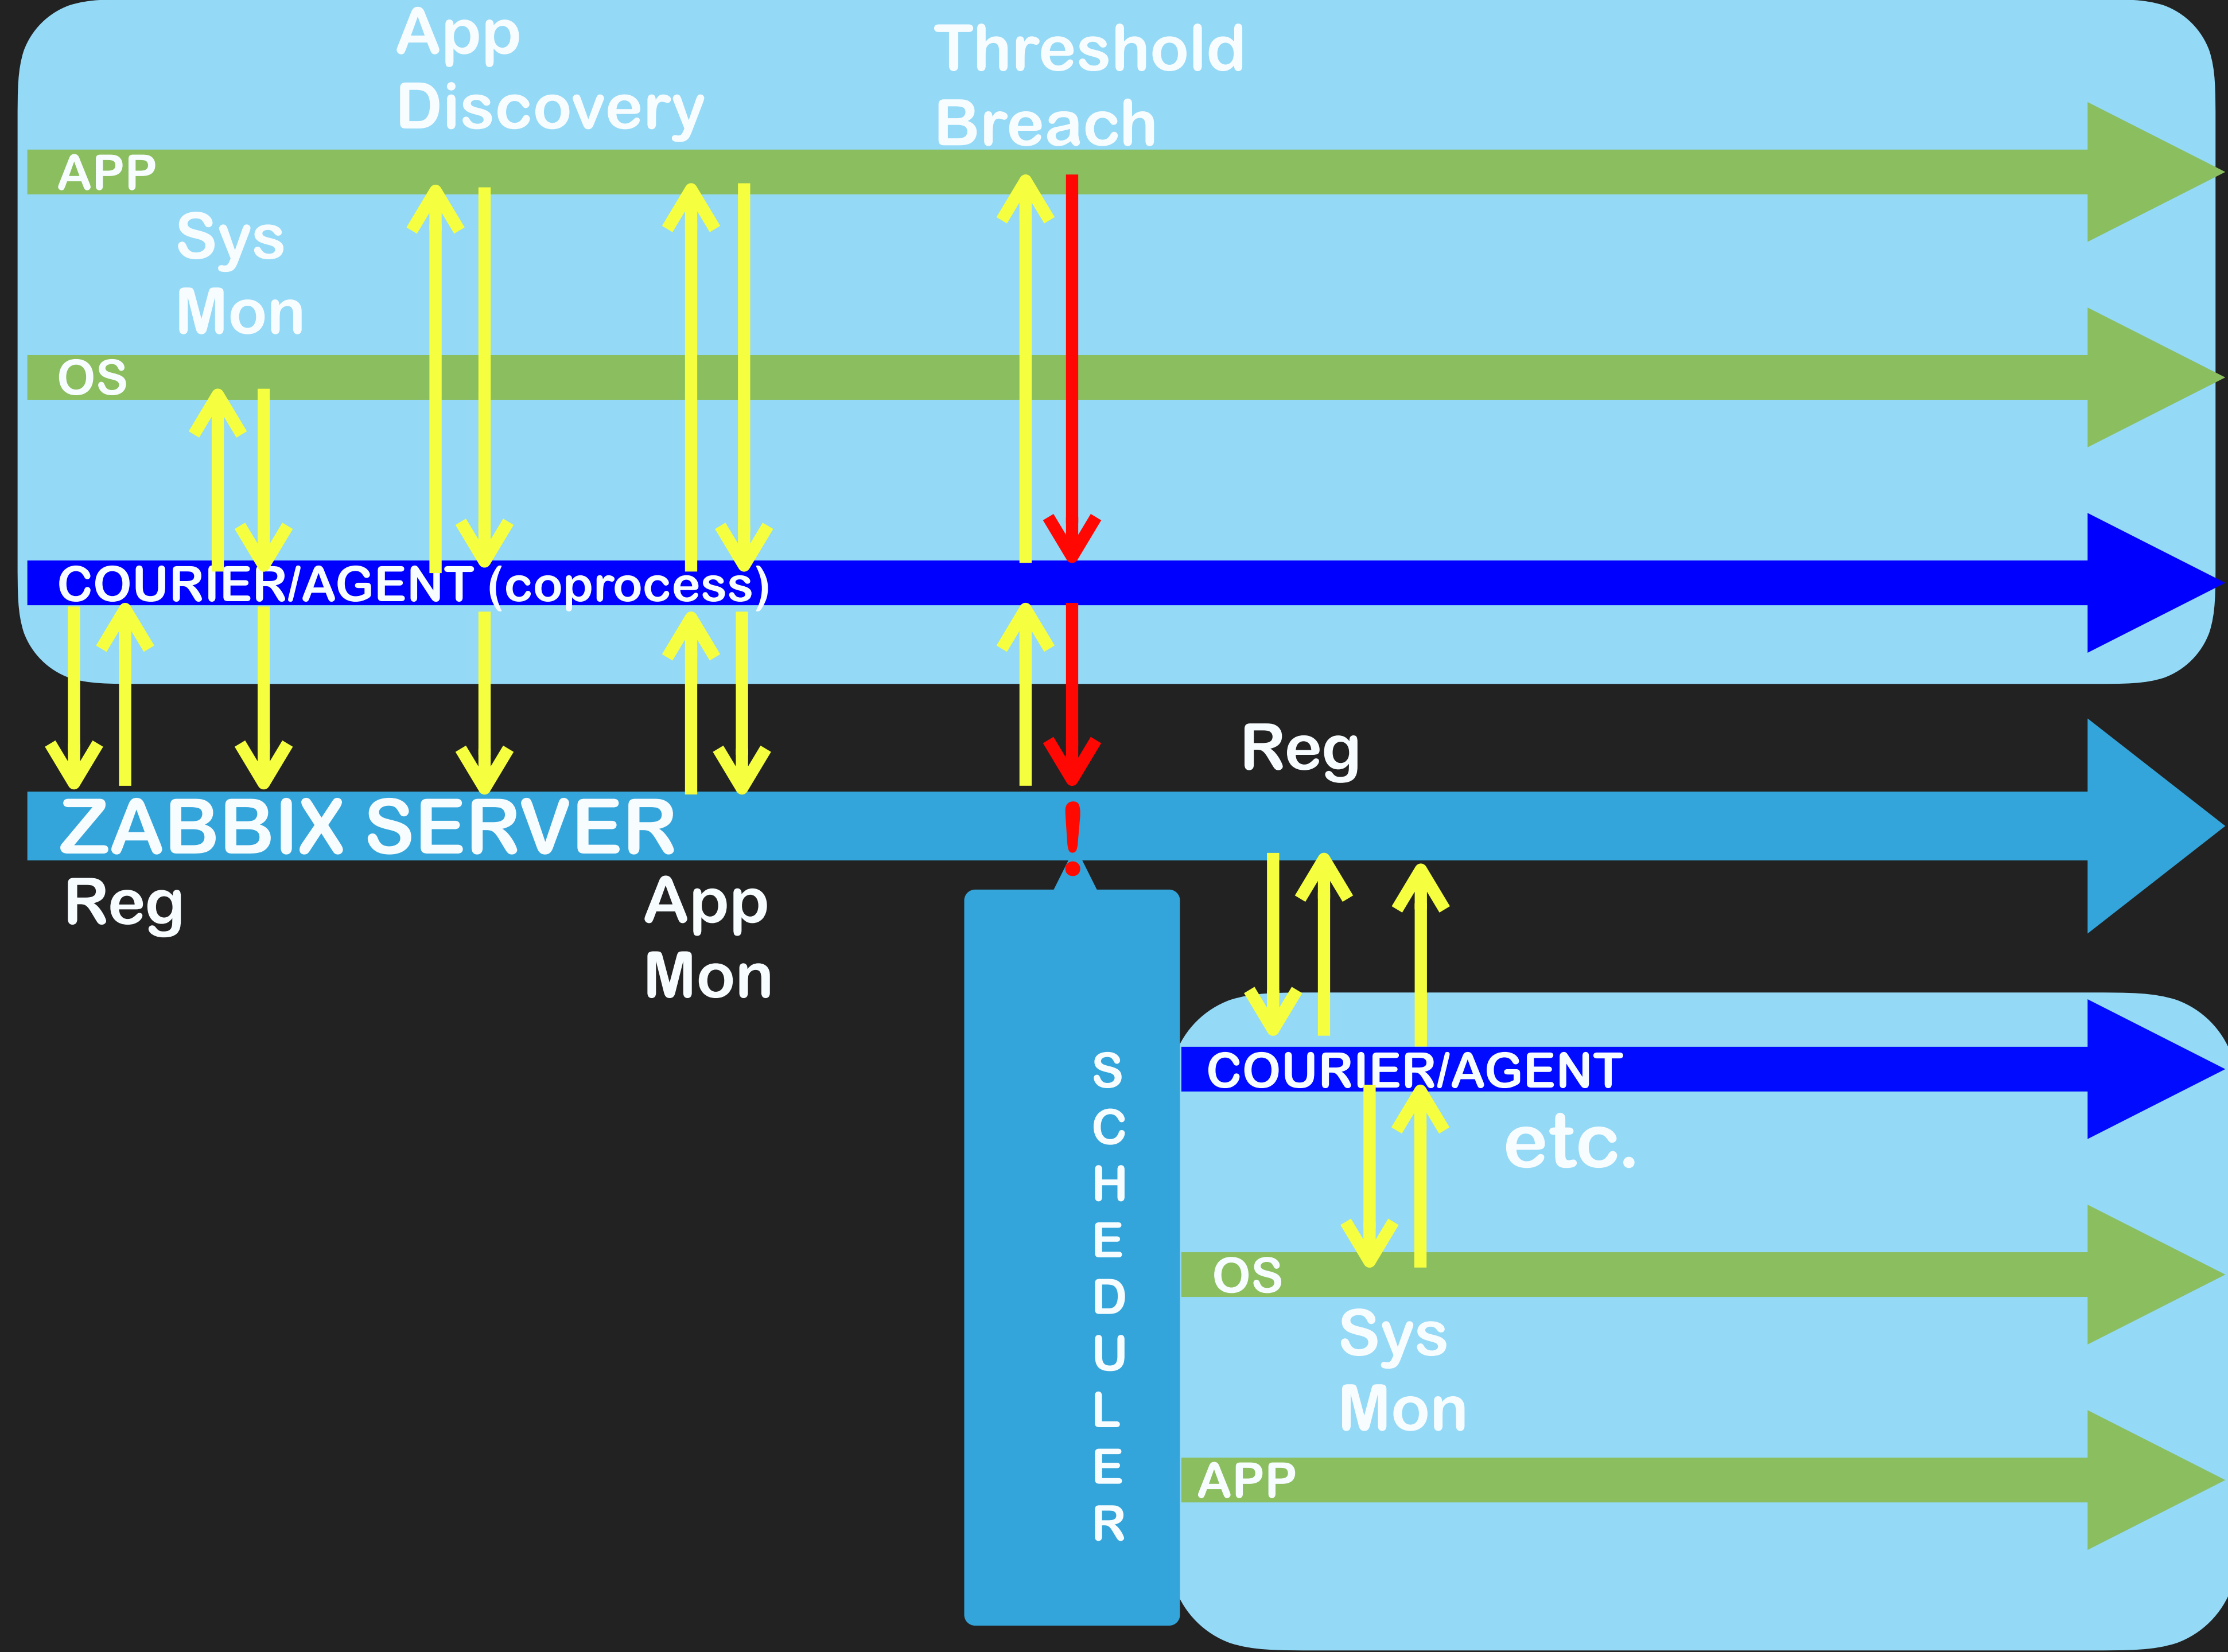
```
# Variable assignment
action=$1; service_name=$2; current_scale=$3; increment=$4

scale_service(){
    /usr/bin/docker-compose --tlsverify --tlscert=${DOCKER_CERT_PATH}cert.pem \
        --tlscacert=${DOCKER_CERT_PATH}ca.pem \
        --tlskey=${DOCKER_CERT_PATH}key.pem --project-name dockerlx \
        --host tcp://dockerapi-private-lab1.mesoform.com:2376 --file /tmp/docker-compose.yml \
        scale ${service_name}=$1
    echo "$(date): Scaled ${service_name} from ${current_scale} to $1" \
        >> /tmp/app_scheduler_output
    exit 0
}

scale_up(){
    desired_scale=$((current_scale + increment))
    scale_service ${desired_scale}
}

scale_down(){
    desired_scale=$((current_scale - increment))
    scale_service ${desired_scale}
}
```











## KEEPING ACTIVE

- ▶ Works for bare metal, VMs and containers
- ▶ application-level guarantees
- ▶ Active versus passive
- ▶ Troubleshooting directly connected services
- ▶ Performance and Reliability

## CONCLUSION


- ▶ Autopilot Pattern and Enfield Method
- ▶ We're already: doing event management, auto-registering, aggregating metrics, performing actions on triggers, maintaining system state, highly optimised, self-healing,
- ▶ Controlling the state
- ▶ Accuracy and performance
- ▶ Short lead time



## WHATS NEXT

- ▶ Load testing
- ▶ Use Zabbix Python interpreter module
- ▶ Key management with vault
- ▶ Swarm/Nomad not docker compose
- ▶ DevOps everything!

## SO LONG AND THANKS FOR ALL THE FISH

- ▶ Read the full article at <http://www.mesoform.com/blog-listing/info/the-concierge-paradigm>
- ▶ Search: "concierge paradigm"
- ▶  @MesoformLtd
- ▶  /mesoform
- ▶  /mesoform
- ▶ <http://www.mesoform.com/contact-us>

